

Load data and initial exploration

```
In [7]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.ensemble import IsolationForest
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
from datetime import datetime

# Load the data
df = pd.read_csv('/content/intrenship_data.csv')
```

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5811 entries, 0 to 5810
Data columns (total 19 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   application_id        5811 non-null   int64
 1   first_name            5811 non-null   object
 2   last_name             5811 non-null   object
 3   email                 5811 non-null   object
 4   phone                 5811 non-null   object
 5   university            5811 non-null   int64
 6   major                 5811 non-null   int64
 7   gpa                   5811 non-null   float64
 8   graduation_year       5811 non-null   int64
 9   submission_date       5811 non-null   object
10  skills                 5811 non-null   object
11  projects              4432 non-null   object
12  ip_address            5811 non-null   object
13  submission_device     5811 non-null   int64
14  time_spent_on_form    5811 non-null   int64
15  is_anomaly            5811 non-null   int64
16  anomaly_type          430 non-null    object
17  submission_hour       5811 non-null   int64
18  submission_day_of_week 5811 non-null   int64
dtypes: float64(1), int64(9), object(9)
memory usage: 862.7+ KB
```

```
In [4]: df.head()
```

Out[4]:

	application_id	first_name	last_name	email	phone	university	major	gpa	graduation_y
--	----------------	------------	-----------	-------	-------	------------	-------	-----	--------------


0	1411	Wyatt	Baker	wyatt.baker79@gmail.com	284.525.2357	61	6	3.45	2
---	------	-------	-------	-------------------------	--------------	----	---	------	---

1	3462	Erika	Ruiz	erika.ruiz43@hotmail.com	602.294.0795	37	5	3.96	2
---	------	-------	------	--------------------------	--------------	----	---	------	---

2	1101	Logan	Davidson	logan.davidson19@hotmail.com	001-264-740-1137x1308	20	8	2.83	2
---	------	-------	----------	------------------------------	-----------------------	----	---	------	---

3	1333	Alicia	Harris	alicia.harris94@yahoo.com	270.319.2564x69677	37	1	3.46	2
---	------	--------	--------	---------------------------	--------------------	----	---	------	---

4	1638	Michael	Bates	michael.bates56@gmail.com	350.722.6922x7910	20	8	3.54	2
---	------	---------	-------	---------------------------	-------------------	----	---	------	---



In [5]: `print(df['is_anomaly'].value_counts())`

```
is_anomaly
0      5381
1       430
Name: count, dtype: int64
```

```
In [6]: print(df['anomaly_type'].value_counts())
```

```
anomaly_type
duplicate      86
bot_pattern    77
rapid_submission 74
inconsistent_data 70
fake_university 63
impossible_gpa  60
Name: count, dtype: int64
```

Data preprocessing

```
In [9]: # Convert submission_date to datetime
df['submission_date'] = pd.to_datetime(df['submission_date'], errors='coerce')
df.dropna(subset=['submission_date'], inplace=True)

# Now proceed with feature extraction
df['submission_hour'] = df['submission_date'].dt.hour
df['submission_minute'] = df['submission_date'].dt.minute
df['submission_day'] = df['submission_date'].dt.day

# Create a feature for submission speed (time spent on form)
df['submission_speed'] = df['time_spent_on_form'] / df['skills'].str.split(',').str.len()

# Create binary features from skills
skills_list = ['Java', 'Python', 'SQL', 'C++', 'Web Development',
               'Machine Learning', 'Data Analysis', 'Cloud Computing',
               'Project Management', 'Communication']
for skill in skills_list:
    df[f'skill_{skill}'] = df['skills'].str.contains(skill).astype(int)

# Handle missing values
df.fillna({'projects': '', 'anomaly_type': 'normal'}, inplace=True)
```

Rule base anomaly detection

```
In [10]: # Detect duplicates based on email, phone, and IP
df['duplicate_email'] = df.duplicated('email', keep=False)
df['duplicate_phone'] = df.duplicated('phone', keep=False)
df['duplicate_ip'] = df.duplicated('ip_address', keep=False)

# Detect rapid submissions (less than 10 seconds)
df['rapid_submission'] = df['time_spent_on_form'] < 10

# Detect impossible GPA values
df['impossible_gpa'] = (df['gpa'] < 0) | (df['gpa'] > 4)

# Detect inconsistent graduation year (before submission year)
df['inconsistent_graduation'] = df['graduation_year'] < df['submission_date'].dt.year

# Detect fake universities (assuming university IDs should be within a certain range)
df['fake_university'] = ~df['university'].between(1, 100)

# Combine rule-based anomalies
df['rule_based_anomaly'] = (df['duplicate_email'] | df['duplicate_phone'] |
                           df['duplicate_ip'] | df['rapid_submission'] |
                           df['impossible_gpa'] | df['inconsistent_graduation'] |
                           df['fake_university'])
```

Machine Learning Anomaly Detection

Isolation Forest

```
In [11]: # Prepare features for Isolation Forest
features = ['gpa', 'time_spent_on_form', 'submission_hour', 'submission_day_of_week'] + \
           [f'skill_{skill}' for skill in skills_list]

X = df[features]

# Train Isolation Forest
```

```
iso_forest = IsolationForest(contamination=0.05, random_state=42)
iso_forest.fit(X)
df['iso_anomaly'] = iso_forest.predict(X)
df['iso_anomaly_score'] = iso_forest.decision_function(X)

# Convert to binary (1 = normal, -1 = anomaly)
df['iso_anomaly'] = df['iso_anomaly'].apply(lambda x: 1 if x == -1 else 0)
```

K-Means Clustering

```
In [12]: # Standardize features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Apply K-Means
kmeans = KMeans(n_clusters=5, random_state=42)
df['cluster'] = kmeans.fit_predict(X_scaled)

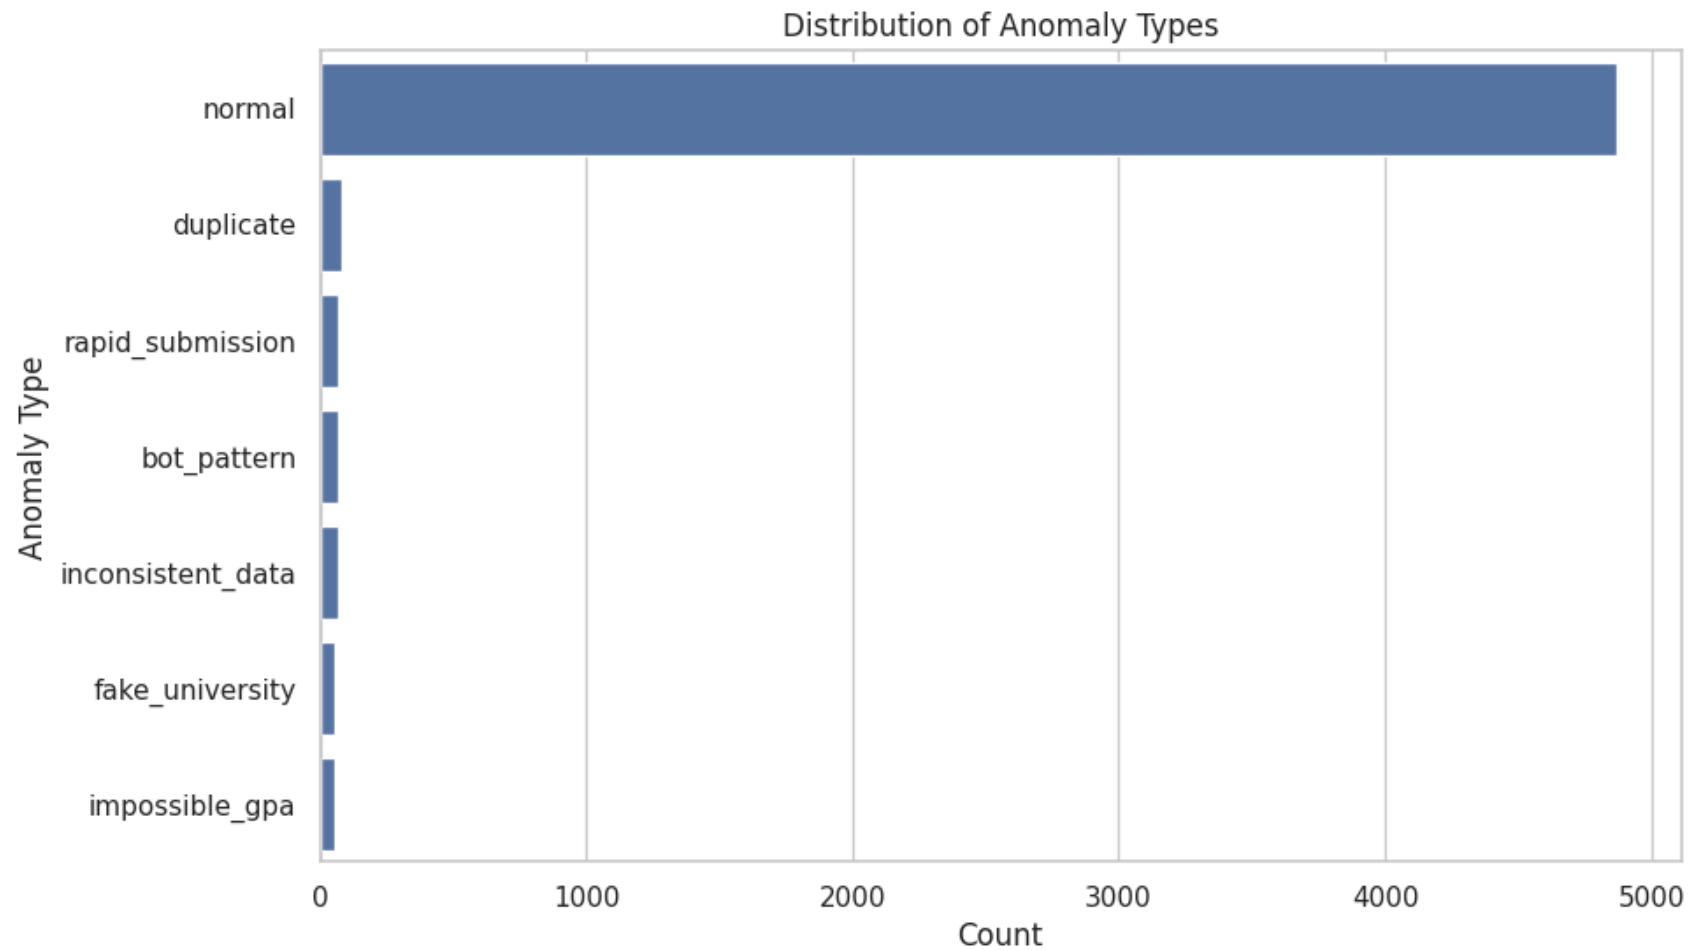
# Calculate distance to cluster centers
distances = kmeans.transform(X_scaled)
df['cluster_distance'] = distances.min(axis=1)

# Flag anomalies as points far from their cluster center
df['kmeans_anomaly'] = (df['cluster_distance'] > df['cluster_distance'].quantile(0.95)).astype(int)
```

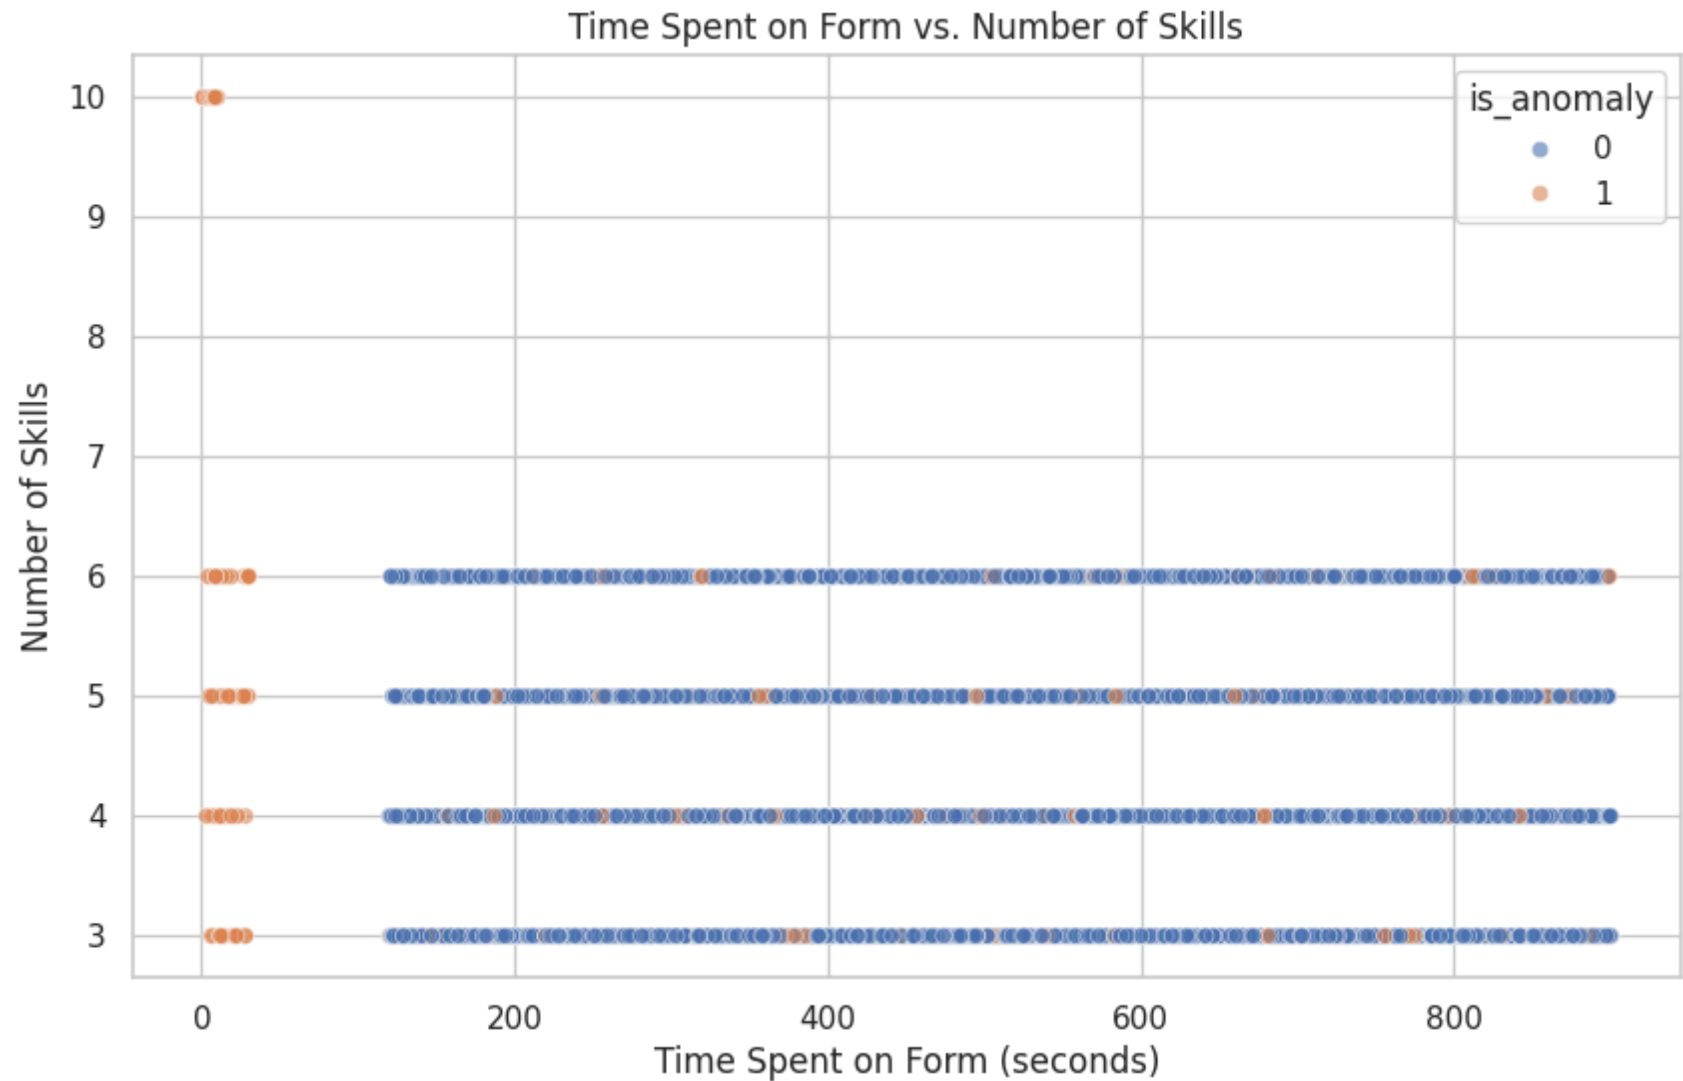
Visualization

```
In [18]: # Set style
sns.set(style="whitegrid")

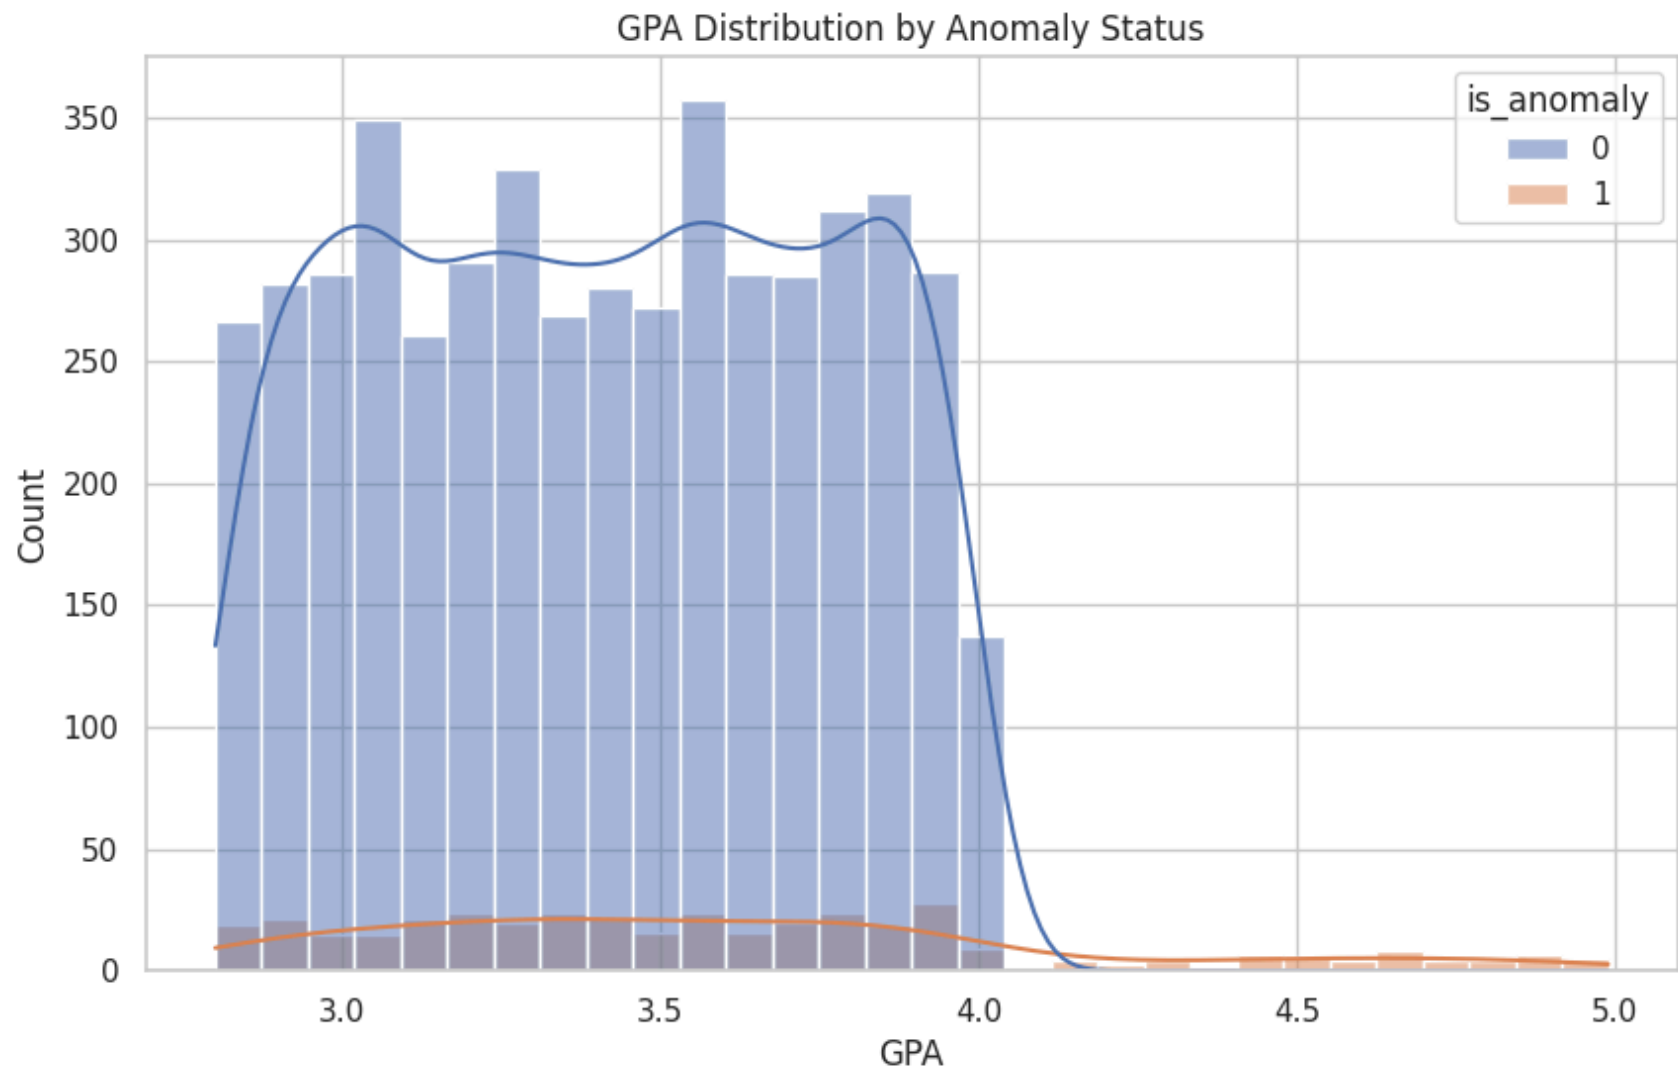
# Plot anomaly types
plt.figure(figsize=(10, 6))
sns.countplot(y='anomaly_type', data=df, order=df['anomaly_type'].value_counts().index)
plt.title('Distribution of Anomaly Types')
plt.xlabel('Count')
plt.ylabel('Anomaly Type')
plt.show()
```




```
In [14]: # Plot time spent on form vs. number of skills
plt.figure(figsize=(10, 6))
sns.scatterplot(x='time_spent_on_form', y=df['skills'].str.split(',').str.len(),
                hue='is_anomaly', data=df, alpha=0.6)
plt.title('Time Spent on Form vs. Number of Skills')
plt.xlabel('Time Spent on Form (seconds)')
plt.ylabel('Number of Skills')
plt.show()
```



```
In [15]: # Plot GPA distribution
plt.figure(figsize=(10, 6))
sns.histplot(data=df, x='gpa', hue='is_anomaly', bins=30, kde=True)
plt.title('GPA Distribution by Anomaly Status')
plt.xlabel('GPA')
plt.ylabel('Count')
plt.show()
```

```
In [16]: # Plot submission hour distribution
plt.figure(figsize=(10, 6))
sns.histplot(data=df, x='submission_hour', hue='is_anomaly', bins=24, kde=True)
plt.title('Submission Hour Distribution by Anomaly Status')
plt.xlabel('Hour of Day')
plt.ylabel('Count')
plt.xticks(range(24))
plt.show()
```




```
In [17]: from sklearn.decomposition import PCA

pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

plt.figure(figsize=(10, 6))
sns.scatterplot(x=X_pca[:, 0], y=X_pca[:, 1], hue=df['is_anomaly'],
               style=df['anomaly_type'], alpha=0.7)
```

```
plt.title('PCA Visualization of Applications with Anomalies Highlighted')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left')
plt.show()
```



6. Alert System Design

```
In [19]: def generate_alerts(row):
          alerts = []

          # Rule-based alerts
          if row['duplicate_email']:
              alerts.append("Duplicate email detected")
```

```

if row['duplicate_phone']:
    alerts.append("Duplicate phone number detected")
if row['duplicate_ip']:
    alerts.append("Duplicate IP address detected")
if row['rapid_submission']:
    alerts.append(f"Extremely rapid submission ({row['time_spent_on_form']} seconds)")
if row['impossible_gpa']:
    alerts.append(f"Impossible GPA value ({row['gpa']})")
if row['inconsistent_graduation']:
    alerts.append(f"Inconsistent graduation year ({row['graduation_year']})")
if row['fake_university']:
    alerts.append(f"Suspicious university ID ({row['university']})")

# ML-based alerts
if row['iso_anomaly'] == 1:
    alerts.append("Isolation Forest anomaly detected")
if row['kmeans_anomaly'] == 1:
    alerts.append("K-Means anomaly detected (far from cluster center)")

return alerts if alerts else ["No alerts"]

# Apply alert generation
df['alerts'] = df.apply(generate_alerts, axis=1)

# Display suspicious applications
suspicious = df[df['alerts'].apply(lambda x: x != ['No alerts'])]
print(f"Found {len(suspicious)} suspicious applications out of {len(df)} total")

# Save suspicious applications to CSV
suspicious.to_csv('suspicious_applications.csv', index=False)

```

Found 2846 suspicious applications out of 5263 total



Key Insights from Visualizations



Anomaly Type Distribution

- **Most Common Anomalies:**
 - `bot_pattern` (repetitive entries)

- rapid_submission
 - fake_university
 - **Less Common but Significant:**
 - impossible_gpa
 - inconsistent_data
-

Time Spent vs. Skills

- Normal applications:
 - Positive correlation between **time spent** and **number of skills**.
 - Anomalous applications:
 - Extremely short times for many skills.
 - Unusually long times for very few skills.
-

GPA Distribution

- Most GPAs are clustered between **2.5 and 4.0**.
 - Anomalies:
 - Values outside the **0-4** range.
 - Some extreme values within the valid range.
-

Submission Times

- Normal applications:
 - Follow a typical daily pattern (peaks during working hours).
 - Anomalous applications:
 - Evenly spread throughout the day.
 - Some concentration at **unusual hours**.
-

Cluster Visualization

- Normal applications:
 - Form **tight clusters** on the PCA plot.
 - Anomalous applications:
 - Appear as **outliers** scattered away from main clusters.
-

Recommendations

Automated Alert System

- Implement **real-time monitoring** to flag applications with multiple anomaly indicators.
- **Prioritize alerts** based on severity:
 - impossible values > duplicates > ML anomalies

Review Process

- Create a **dashboard** for human reviewers:
 - Display flagged applications with full details.
 - Show anomaly scores for transparency.

Continuous Improvement

- **Update ML models** regularly with new data and confirmed fraud cases.
- **Adjust rule thresholds** based on historical patterns and trends.

Preventive Measures

- Implement **CAPTCHA** or other bot-prevention mechanisms.

- **Validate university IDs** against a known list.
 - **Rate-limit submissions** from the same IP address.
-

This **hybrid approach** (combining rule-based detection and ML) ensures **effective identification** of suspicious internship applications while minimizing **false positives**.

In []: